

# Assignment 6

COMPSCI 2CO3: Data Structures and Algorithms–Fall 2025

Deadline: November 14, 2025

Department of Computing and Software  
McMaster University

Please read the *Course Outline* for the general policies related to assignments.

**Plagiarism is a serious academic offense and will be handled accordingly.**

**All suspicions will be reported to the Office of Academic Integrity  
(in accordance with the Academic Integrity Policy).**

This assignment is an *individual* assignment: do not submit work of others. All parts of your submission *must* be your own work and be based on your own ideas and conclusions. Only *discuss or share* any parts of your submissions with your TA or instructor. You are *responsible for protecting* your work: you are strongly advised to password-protect and lock your electronic devices (e.g., laptop) and to not share your logins with partners or friends!

If you *submit* work, then you are certifying that you are aware of the *Plagiarism and Academic Dishonesty* policy of this course outlined in this section, that you are aware of the *Academic Integrity Policy*, and that you have completed the submitted work entirely yourself. Furthermore, by submitting work, you agree to automated and manual plagiarism checking of all submitted work.

*Late submission policy.* We provide a general *shared grace period* for all assignments of 5 days (120h) that students can utilize whichever way they seem fit without further notice or approval. For example, a student can hand in *one* assignment 100h late, or hand in 4 assignments 20h late (as long as the total stays below 120h). Beyond the above grace period, we require students to *submit on time*. Late submissions are not accepted and will result in a grade of zero. The above grace period *does not stack* with MSAFs or SAS accommodations.

**Problem 1.** Consider the sequence of values  $S = [78, 6, 88, 19, 77, 79, 94, 7, 1, 36, 934]$ .

P1.1. Draw the *left-leaning* red-black tree obtained by adding the values in  $S$  in sequence. Show each step.

P1.2. Consider the hash function  $h(x) = (x + 5) \bmod 12$  a hash-table of 12 table entries that uses hashing with separate chaining. Draw the hash-table obtained by adding the values in  $S$  in sequence. Show each step.

P1.3. Consider the hash function  $h(x) = (x + 5) \bmod 12$  a hash-table of 12 table entries that uses hashing with linear probing. Draw the hash-table obtained by adding the values in  $S$  in sequence. Show each step.

Do *not* spend time drawing beautiful trees or tables: a clear textual representation is good enough.

**Problem 2.** Consider a list  $L$  of pairs. A common operation on such a list is *aggregating* the list on a given field and then counting, per value of that field, the number of tuples that have that value. For example, consider a list of pairs that holds course enrollment information (pairs *student identifier* and *course code*). If we aggregate this list on *student identifier*, then we compute, per *student*, the number of courses that student takes. Likewise, if we aggregate this list on *course code*, then we compute, per *course*, the number of students enrolled in that course.

P2.1. Write an algorithm that, given a set of pairs  $L$ , computes the aggregate

$$\text{aggr}(L) = \{(a, \text{count}(a, L)) \mid \text{count}(a, L) \neq 0\} \text{ with } \text{count}(a, L) = |\{b \mid (a, b) \in L\}|.$$

The complexity of your algorithm must be expected  $O(|L|)$ , your algorithm cannot change the list  $L$ , and your algorithm can use at-most  $O(|\text{result}|)$  memory with  $|\text{result}|$  the size of the result (the number of pairs  $(a, \text{count}(a, L)) \in \text{aggr}(L)$ ).

You can only assume that you can efficiently *iterate* over all values in  $L$ . You *cannot* make any further assumptions on how list  $L$  is structured (e.g., you cannot assume that  $L$  is sorted).

**HINT:** Feel free to first solve this problem when you know that  $|\text{result}| = k$  beforehand. Next, argue why you do not need this assumption in practice.

P2.2. Argue why your algorithm correctly computes  $\text{aggr}(L)$  and argue why your algorithm meets the complexity requirements stated in Problem P2.1.

Consider two sets-of-pairs  $F$  and  $S$ . For example, set  $F$  can hold student information (*student identifier* and *name*) and list  $S$  can hold course enrollment information (*student identifier* and *course identifier*). When processing data, a typical operation on  $F$  and  $S$  is computing the *semi-join*:

$$\text{SemiJoin}(F, S, \delta) = \{(a, b) \mid (a, b) \in F \wedge (a, \delta) \in S\}.$$

e.g., if  $F$  holds student information,  $S$  holds course enrollment information, and  $\delta = \text{'COMPSCI 2C03'}$ , then this semi-join will return the student information of all students that take the course COMPSCI 2C03.

P2.3. Write an algorithm that, given sets-of-pairs  $F$  and  $S$  and value  $\delta$ , *correctly computes* the above semi-join  $\text{SemiJoin}(F, S, \delta)$ . The complexity of your algorithm must be expected  $O(|F| + |S|)$ , your algorithm cannot change the lists  $F$  and  $S$ , and your algorithm can use at-most  $O(|S|)$  extra memory besides the memory used to store the result.

You can only assume that you can efficiently *iterate* over all values in  $F$  and  $S$ . You *cannot* make any further assumptions on how  $F$  and  $S$  are structured (e.g., you cannot assume that they are sorted).

**HINT:** Feel free to first solve this problem when you know that  $|\text{result}| = k$  beforehand. Next, argue why you do not need this assumption in practice.

P2.4. Argue why your algorithm correctly computes  $\text{SemiJoin}(F, S, \delta)$  and argue why your algorithm meets the complexity requirements stated in Problem P2.3.

## Assignment Details

Write a report in which you solve each of the above problems. Your submission:

1. must start with your name, student number, and MacID;
2. must be a PDF file;
3. must have clearly labeled solutions to each of the stated problems;
4. must be clearly presented;
5. must *not* be hand-written: prepare your report in L<sup>A</sup>T<sub>E</sub>X or in a word processor such as Microsoft Word (that can print or export to PDF).

**Submissions that do not follow the above requirements will get a grade of zero.**

## Grading

Each problem counts equally toward the final grade of this assignment.